

Customer(customer-name,street,city)

Branch(branch-name,account)

Account(customer-name,branch-name,account-number)

(a)Find the names of all customers who have an account in the 'Region12' branch→

$\pi_{\text{customer-name}}(\sigma_{\text{branch-name}='Region 12'}(\text{Account}))$

(b)Find the names of all customers who have an account in a branch NOT located in the same city that they live in→

$\pi_{\text{customer-name}}(\sigma_{\text{A.city} \neq \text{B.city} \wedge \text{A.branch-name}=\text{B.branch-name}}(\rho_B(\text{Branch}) \times \rho_A(\text{Customer} \bowtie \text{Account})))$

(c)Find the branches that do not have any accounts→

$\pi_{\text{branch-name}}(\text{Branch}) - \pi_{\text{branch-name}}(\text{Account})$

(d)Find the customer names who do not have any account in the 'Region12' branch→

$\pi_{\text{customer-name}}(\text{Customer}) - \pi_{\text{customer-name}}(\sigma_{\text{branch-name}='Region 12'}(\text{Account}))$

(e)Find the customer names who have accounts in all the branches located in 'Los Angeles' →

$\pi_{\text{customer-name}}(\text{Customer}) - \pi_{\text{customer-name}}(\pi_{\text{customer-name}}(\text{Customer}) \times \pi_{\text{branch-name}}(\sigma_{\text{city}='Los Angeles'}(\text{Branch}))) - \pi_{\text{customer-name,branch-name}}(\text{Account})$

(f) Find the customer names who have only one account →

$\pi_{\text{customer-name}}(\text{Customer}) - \pi_{\text{A.customer-name}}(\sigma_{\text{A.branch-name} \neq \text{B.branch-name} \vee \text{A.account-number} \neq \text{B.account-number} \wedge \text{A.customer-name}=\text{B.customer-name}}(\rho_A(\text{Account}) \times \rho_B(\text{Account})))$

Student(sid,GPA)

Find the ids of the students with the lowest GPA

$\pi_{\text{sid}}(\text{Student}) - \pi_{\text{A.sid}}(\sigma_{\text{A.GPA} > \text{B.GPA} \wedge \text{A.sid} \neq \text{B.sid}}(\rho_A(\text{Student}) \times \rho_B(\text{Student})))$

Employee(person-name,age,street,city)

Work(person-name,company-name,salary)

Company(company-name,city)

Manage(person-name,manager-name)

(a)Find the name(s) of the employee(s) whose total salary is higher than those of all employees living in Los Angeles →SELECT person-name FROM Work GROUP BY person-name HAVING SUM(salary) >ALL (SELECT SUM(salary) FROM Work W, Employee E WHERE W.person-name=E.person-name AND city='Los Angeles' GROUP BY W.person-name);

(b)Find the names of the manager(s) whose total salary is higher than that of at least one employee that they manage → SELECT person-name FROM Work W1 GROUP BY person-name HAVING SUM(salary) >SOME (SELECT SUM(salary) FROM Work W2, Manage M WHERE W1.person-name=M.manager-name AND M.person-name=W2.person-name GROUP BY W2.person-name)

ComputerProduct(manufacturer,model,price)

Desktop(model,speed,ram,hdd)

Laptop(model,speed,ram,hdd,weight)

(a)Using two INSERT statements, insert a desktop computer manufactured by HP, with model number 1200, price \$1000, speed 1.2Ghz, 256 MB RAM, and an 80GB hard drive → INSERT INTO ComputerProduct VALUES ('HP',1200,1000); INSERT INTO Desktop VALUES (1200, 1.2, 256, 80);

(b)Using two DELETE statements, delete all desktops manufactured by IBM with price below \$1000 → DELETE FROM Desktop WHERE model IN (SELECT model FROM ComputerProduct WHERE manufacturer='IBM' AND price <

Data Integrity Enforcement

PRIMARY KEY(dept,cnum,sec)

UNIQUE(dept,sec,title)

FOREIGN KEY sid REFERENCES Student(sid)

ON DELETE/UPDATE CASCADE/SET NULL

Disks

Seek time avg: 10msec

Q:What is the overall capacity? → KB/disk (1,000,000,000KB=1TB)

Access Time=(seek time)+(rotational delay)+(transfer time)

Q:For 6000 RPM, average rotational delay? →

100 rotations/sec, 0.01sec/rotation=10msec. average is half, so 5msec

Q:6000RPM,10000sectors/track. How long to read one sector? →

10 msec/10000=0.001msec

Q:6000RPM,10000sectors/track. Average access time to read one sector? →

(seek time)(10ms)+(rotational delay)(5ms)+(transfer time)(0.001)=15.001ms

Q:6000RPM,10000sectors/track,1KB/sector →

What is the transfer rate? →Burst Transfer rate=(RPM/60)*(sectors/track)*(bytes/sector)

CHECK Constraints

Q1:Class snum should be <600 and class units should be <10 → CHECK(cnum<600 AND unit<10)

Q2:The units of all CS classes should be >3 → CHECK(dept<>'CS' OR unit>3)

Q3:Students whose GPA<2 cannot take CS class → CHECK((sid IN(SELECT sid FROM Student WHERE GPA>=2)) OR dept<>'CS')

Q4:Can we express referential integrity constraint using CHECK constraint → CHECK(sid IN(SELECT sid FROM Student));

TRIGGER Constraints

CREATE TRIGGER <name>

AFTER/BEFORE INSERT/UPDATE/DELETE ON R

REFERENCING NEW/OLD ROW/TABLE AS <referencing clause>

FOR EACH ROW/STATEMENT

WHEN(<condition>)

BEGIN

<event>

END

Sequential vs. Random I/O

Q: How long to read 3 sequential sectors? →

10ms*(3/10000)=0.003ms

Q: How long to read 3 random sectors? (head above 1st scctr→ 10ms/10000=0.001ms+(10ms+5ms+0.001ms)+(10ms+5ms+0.001ms)=30.003ms

Equal?

1. $\pi_A(R - S) = \pi_A(R) - \pi_A(S) \rightarrow \text{NO}; (1,2), (1,2)$
SELECT B FROM R GROUP BY B; SELECT DISTINCT E FROM R → YES
SELECT B FROM R WHERE NOT EXISTS (SELECT * F WHERE R.B=S.B); (SELECT B FROM R) EXCEPT (SEL FROM S) → NO, R=(1,0),(2,); S={}
SELECT B FROM R R1 WHERE R <= ALL (SELECT B F R R2 WHERE R1.A<>R2.A)); SELECT MIN(B) FROM I R=(1,0),(2,0)

2. Manager: CHECK (salary > ALL (SELECT salary FR Employee WHERE E.dept=M.dept))
Employee: CHECK (salary < ALL (SELECT salary FRO Manager WHERE M.dept = E.dept));

3. 800GB=10 surfaces x 10,000 tracks x 8GB x X sec/
X = 1000 sec/track
10ms + 5ms + 10ms/1000 = 15.01 ms

4. 40B/tuple; 1,000,000 tuples. Blocks (8GB) needed =1,000,000/200tuple/block=5,000 blocks

Overall capacity: 2 platters x 2 surfaces/platter, 2000 tracks/surface x 1KB/sector=1,600,000,000 KB/disk TB/disk

Terminology:

- Relations=tables
- Attributes=columns
- Tuples=rows
- Domain=type
- Data model=graph/tree model
- Schema:structure of relations(variable type)
- EX:Student(sid:int,name:varchar(9),addr:varchar(9))
- Instance=data(value)
- Keys:set of attributes that uniquely identifies a tuple
- EX:Student(sid,name,addr)
- Set semantics:no duplicate tuples(relational algebra)
- Multi-set semantics:duplicate

Tables:

- One primary key per table
- UNIQUE for other keys
- SQL92:no NULL in primary
- DEFAULT for default

5 Steps Database Construction:

1. Domain analysis
2. Database design(E/R model, database design theory)
3. Table creation(DDL)
4. Load

Query and update(DML)

SQL Data Types:

String
Char(n)—fixed length
Varchar(n)—variable length
Number
Integer—32 bit
Decimal (5,2)—999.99
Real, double—32 bit,64 bit
Datetime
Date—'2010-01-15'
Time—'13:50:00'
Timestamp—'2010-01-15 13:50:00'

SQL Table Creation/Deletion:

```
CREATE TABLE Course (  
    dept CHAR(2) NOT NULL  
    DEFAULT 'CS',  
    cnum INT NOT NULL,  
    sec INT NOT NULL,  
    title VARCHAR(50),  
    PRIMARY KEY(dept,cnum,sec),  
    UNIQUE(dept,sec,title)  
);  
DROP TABLE Course;
```

Loading Data (MySQL): LOAD
DATA INFILE <datafile> INTO
TABLE <table>

Relational Algebra:

SELECT (σ)
PROJECT (π)
CROSS PRODUCT (\times)
NATURAL JOIN (\bowtie)
RENAME (ρ)
UNION (\cup)
SET DIFFERENCE ($-$)
INTERSECT (\cap)
DIVISION ($/$)

EX:

Student(sid,name,addr,age,GPA)
Class(dept,cnum,sec,unit,title,instructor)
Enroll(sid,dept,cnum,sec)

Q1:All students→Student

Q2:Students with age<18→ $\sigma_{\text{age}<18}(\text{Student})$

Q3:Students with GPA>3.7 and age<18→ $\sigma_{\text{GPA}>3.7 \wedge \text{age}<18}(\text{Student})$

Q4:sid and GPA of all students→ $\pi_{\text{sid,GPA}}(\text{Student})$

Q5:All departments offering a class→ $\pi_{\text{dept}}(\text{Class})$

Q6:sid and GPA of students with age<18→ $\pi_{\text{sid,GPA}}(\sigma_{\text{age}<18}(\text{Student}))$

Q7:Names of students who take CS classes

→ $\pi_{\text{name}}(\sigma_{\text{Student.sid=Enroll.sid}}(\text{Student} \times (\sigma_{\text{dept='CS'}}(\text{Enroll}))))$

Q8:Names of students who take CS classes→

$\pi_{\text{name}}(\text{Student} \bowtie (\sigma_{\text{dept='CS'}}(\text{Enroll})))$

Q9:Names of students who take classes from "John Cho"→

$\pi_{\text{name}}(\text{Student} \bowtie (\text{Enroll} \bowtie (\sigma_{\text{instructor='John Cho'}}(\text{Class}))))$

Q10:Names of the student pairs who live in the same address→

$\pi_{s1.name=s2.name}(\sigma_{s1.addr=s2.addr \wedge s1.sid=s2.sid}(\rho_{s1}(\text{Student}) \times \rho_{s2}(\text{Student})))$

Q11: All people's names→ $\pi_{\text{name}}(\text{Student}) \cup \rho_{c(\text{name})}(\pi_{\text{instructor}}(\text{Class}))$

Q12:Courses(dept,cnum,sec) that no one takes→ $\pi_{\text{dept,cnum,sec}}(\text{Class}) -$

$\pi_{\text{dept,cnum,sec}}(\text{Enroll})$

Q14:Instructor names who teach both CS and EE courses→

$\pi_{\text{instructor}}(\sigma_{\text{dept='CS'}}(\text{Class})) \cap \pi_{\text{instructor}}(\sigma_{\text{dept='EE'}}(\text{Class}))$

Q15:Sids of students who take every CS class→

AllSids— $\pi_{\text{sid}}(\text{AllSids} \times \text{CSClasses} - \text{Enroll})$

$\pi_A(R) - \pi_A(\pi_A(R) \times S - R) = R/S$

R←enroll	
A (Sid)	B (ClassID)
a ₁	b ₁
a ₁	b ₂
a ₂	b ₁
a ₃	b ₂

A
a ₁

S←CS Classes
B
b ₁
b ₂

MySQL:

DISTINCT (to remove duplicates)

% = any string

_ = one character

Set operators:

INTERSECT, UNION, EXCEPT

- same schema for operands

- based on set semantics

- keep duplicates with 'ALL'

Subqueries:

- can rewrite them as long as we don't have negation

- with negation, we need EXCEPT

Set membership (IN, NOT IN)

- a IN R is TRUE is a is in R

Set comparison (> ALL, < SOME,)

- <> ALL \equiv NOT IN, = SOME \equiv IN

Set operators ($\cup, \cap, -$)

- NULL is treated like other values here

- check NULL: IS NULL, IS NOT NULL

Arithmetic op's and comp's

- for NULL's, UNKNOWN, only true values returned

Three-valued logic:

Truth table:

- AND: $U \ \& \ T = U, U \ \& \ F = F, U \ \& \ U = U$

- OR: $U \ | \ T = T, U \ | \ F = U, U \ | \ U = U$

LEFT/RIGHT/FULL OUTER JOIN

Data Modification

INSERT INTO Enroll VALUES (301,'CS',201,01);

DELETE FROM Student WHERE sid NOT IN

(SELECT sid FROM Enroll);

UPDATE Class SET cnum=cnum+100 WHERE dept='CS';

Q1:Titles and instructors of all CS classes→

SELECT title, instructor FROM Class WHERE dept='CS';

Q2:Names and GPAs of all students who take CS class(es)→

SELECT DISTINCT name, GPA AS Grade, FROM Student S, Enroll E WHERE dept='CS' AND S.sid=E.sid;

Q3:All student names and GPAs who live on Wilshire→

SELECT name,GPA FROM Student WHERE addr LIKE '%Wilshire%';

Q4:People's names both students and instructors→

(SELECT name FROM Student) UNION (SELECT instructor-name FROM Class);

Q5:Sids of students who do not take any CS class→

(SELECT sid FROM Student) EXCEPT (SELECT sid FROM Enroll WHERE dept='CS');

Q6:Sides who live with student 301→

SELECT sid FROM Student WHERE addr=(SELECT addr FROM Student WHERE sid=301);

Q7:Student names who take CS classes→

SELECT name FROM Student S WHERE sid IN (Select sid FROM Enroll WHERE dept='CS');

SELECT name FROM Student S, Enroll E WHERE S.sid=E.sid AND dept='CS';

Q8:Students names who take no CS class→

SELECT name FROM Student S WHERE sid NOT IN (SELECT sid FROM Enroll WHERE dept='CS');

(SELECT name FROM Student) EXCEPT (SELECT name FROM Enroll E, Student S WHERE dept='CS' AND E.sid=S.sid);

Q9: Student IDs who has higher GPA than any student of age 18 or less→

SELECT sid FROM Student WHERE GPA>SOME(SELECT GPA FROM Student WHERE age<=18);

Q10:Student IDs whose GPA is higher than at least one student of age 18 or less→

SELECT sid FROM Student WHERE GPA>SOME(SELECT GPA FROM Student WHERE age<=18);

Q11:Student names who take any class→SELECT name FROM Student S WHERE EXI (SELECT * FROM Enroll E WHERE S.sid=E.sid);

Q12:Average GPA of all students→SELECT AVG(GPA) FROM Student;

Q13:Number of students taking CS classes→

SELECT COUNT(DISTINCT sid) FROM Enroll WHERE dept='CS';

Q14:Average GPA of students who take CS classes→

SELECT AVG(GPA) FROM Student S, Enroll E WHERE dept='CS' AND S.sid=E.sid;

SELECT AVG(GPA) FROM Student S WHERE sid IN (SELECT sid FROM Enroll E WHERE dept='CS');

Q15:Average GPA for each group→

SELECT age, AVG(GPA) FROM Student GROUP BY age;

Q16:Number of classes each student takes→

SELECT sid, COUNT(*) FROM Enroll GROUP BY sid;

Q16 Using Outer Join: SELECT sid, COUNT(*) FROM Student S LEFT OUTER JOIN Enr E ON S.sid=E.sid GROUP BY sid;

Q17:Students who take two classes or more→

SELECT sid FROM Enroll GROUP BY sid HAVING COUNT(*)>=2;